

IA PUCP - Diplomado de Desarrollo de Aplicaciones de Inteligencia Artificial
Python para Ciencia de Datos



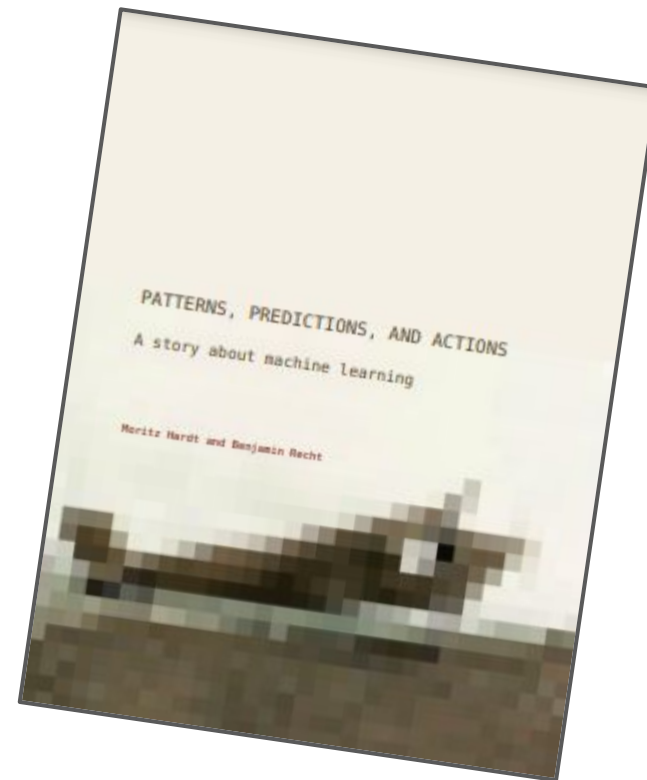
Introducción a Frameworks de Deep Learning y Redes Neuronales

Lectura Recomendada


Patterns, predictions, and actions: A story about machine learning

Moritz Hardt and Benjamin Recht


<https://mlstory.org/>





Reconocimiento de patrones





Reconocimiento de patrones



Reconocimiento de patrones



Reconocimiento de patrones



Machine Learning

¿Qué es Machine Learning?

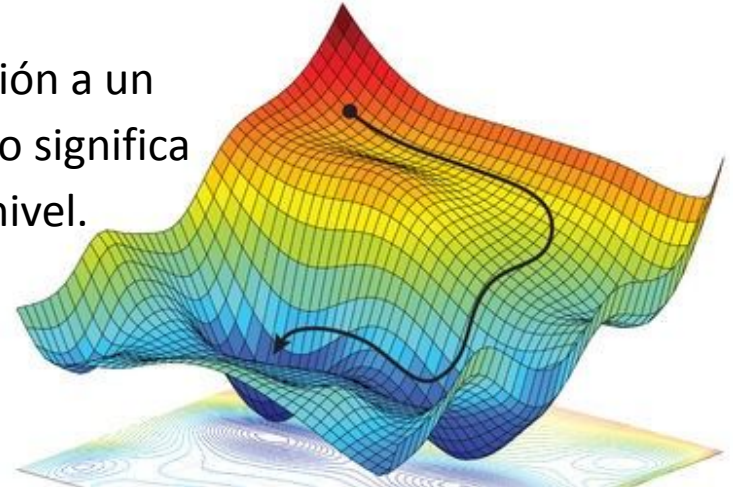
“Se puede decir que un programa aprende de una experiencia **E** con respecto a cierto tipo de tareas **T** con una medida de performance **P** si la performance en las tareas del tipo **T**, medidas por **P**, mejora con la experiencia **E**.”

Tom Mitchell

Convertir un problema de aprendizaje en uno de optimización

https://www.sciencemag.org/sites/default/files/styles/inline_699w_no_aspect/public/ma_0504_NID_alchemy_WEB.jpg?itok=YufMqkHl

- No debemos olvidar que los modelos de machine learning son construidos al optimizar un objetivo **(usualmente minimizar una función de pérdida)**
- Existen múltiples formas en las que un objetivo puede satisfacerse
- Podemos medir empíricamente la generalización a un conjunto de evaluación. Incluso si sale bien, no significa que se estén entendiendo conceptos de alto nivel.



Aproximaciones a la clasificación de imágenes

<https://www.slideshare.net/yannex/yann-le-cun> (2013)



Traditional Pattern Recognition: Fixed/Handcrafted Feature Extractor



Mainstream Modern Pattern Recognition: Unsupervised mid-level features



Deep Learning: Representations are hierarchical and trained



Redes Neuronales: Funciones Parametrizables



...o programas sub-especificados

Las redes neuronales pueden ser vistas como “familias de funciones”

El funcionamiento va a depender de los parámetros que escojamos




Modelo de Neurona McCulloch & Pitts (1943)



<https://upload.wikimedia.org/wikipedia/commons/4/44/Neuron3.png>

https://www.researchgate.net/publication/323465059/figure/fig2/AS:599207769554946@1519873673906/McCulloch-Pitts-computational-model-of-a-neuron_W640.jpg

Perceptrons



Division of Rare and Manuscript Collections. Frank Rosenblatt '50, Ph.D. '56, works on the “perceptron” – what he described as the first machine “capable of having an original idea.”

<https://news.cornell.edu/stories/2019/09/professors-perceptron-paved-way-ai-60-years-too-soon> 12

Perceptrons



“Yet we are about to witness the birth of such a machine – a machine capable of perceiving, recognizing and identifying its surroundings without any human training or control.” (1958)

Division of Rare and Manuscript Collections. Frank Rosenblatt '50, Ph.D. '56, works on the “perceptron” – what he described as the first machine “capable of having an original idea.”

<https://news.cornell.edu/stories/2019/09/professors-perceptron-paved-way-ai-60-years-too-soon> 13

Perceptrons

Division of Rare and Manuscript Collections. The first page of Rosenblatt's article, "The Design of an Intelligent Automaton," in Research Trends, a Cornell Aeronautical Laboratory publication, Summer 1958.

<https://news.cornell.edu/stories/2019/09/professors-perceptron-paved-way-ai-60-years-too-soon>



Perceptrons (1958)

¹ From the New York Times archives:

NEW NAVY DEVICE LEARNS BY DOING

Psychologist Shows Embryo
of Computer Designed to
Read and Grow Wiser

WASHINGTON, July 7 (UPI)
—The Navy revealed the embryo of an electronic computer today that it expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence.

The embryo—the Weather Bureau's \$2,000,000 "704" computer—learned to differentiate between right and left after fifty attempts in the Navy's demonstration for newsmen.

The service said it would use this principle to build the first of its Perceptron thinking machines that will be able to read and write. It is expected to be finished in about a year at a cost of \$100,000.

<https://mlstory.org/>

<https://www.nytimes.com/1958/07/08/archives/new-navy-device-learns-by-doing-psychologist-shows-embryo-of.html>


Perceptrons



*“...be able to walk,
talk, see, write,
reproduce itself and
be conscious of its
existence.”*

Perceptron: Separador Lineal

- En el caso de clasificación binaria, consideremos las clases $\{-1, 1\}$
- El algoritmo perceptrón busca encontrar un separador lineal de la data (o un hiperplano)




<https://mlstory.org/>


Perceptron: Separador Lineal

Division of Rare and Manuscript Collections.
An image of the perceptron from Rosenblatt's "The Design of an Intelligent Automaton," Summer 1958.


<https://news.cornell.edu/stories/2019/09/professors-perceptron-paved-way-ai-60-years-too-soon>



Minsky & Papert y el AI-Winter (1969)



OR




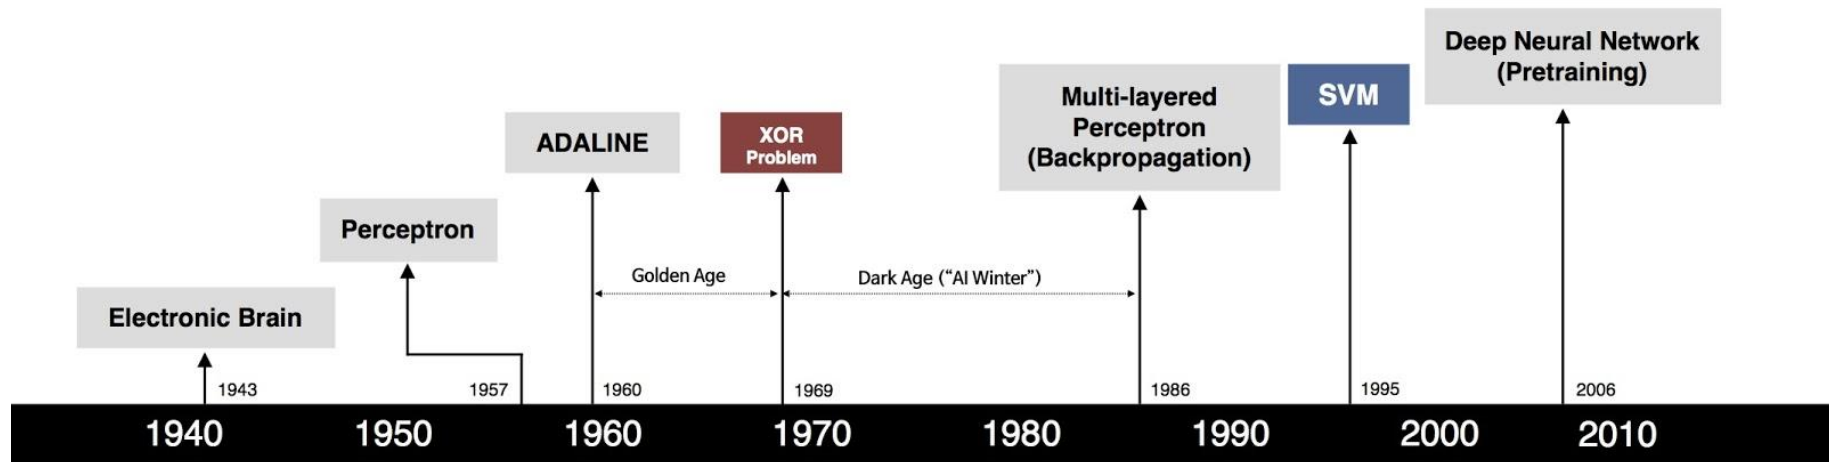
XOR

The XOR problem


Input 1	Input 2	Output
0	0	0
0	1	1
1	1	0
1	0	1

<https://dev.to/jbahire/de-mystifying-the-xor-problem-1blk>





S. McCulloch – W. Pitts




- Adjustable Weights
- Weights are not Learned



F. Rosenblatt




B. Widrow – M. Hoff



- Learnable Weights and Threshold




M. Minsky – S. Papert



- XOR Problem




D. Rumelhart – G. Hinton – R. Williams



- Solution to nonlinearly separable problems
- Big computation, local optima and overfitting




V. Vapnik – C. Cortes



- Limitations of learning prior knowledge
- Kernel function: Human Intervention



G. Hinton – S. Ruslan




- Hierarchical feature Learning


Hubel & Wiesel (1959-1962)



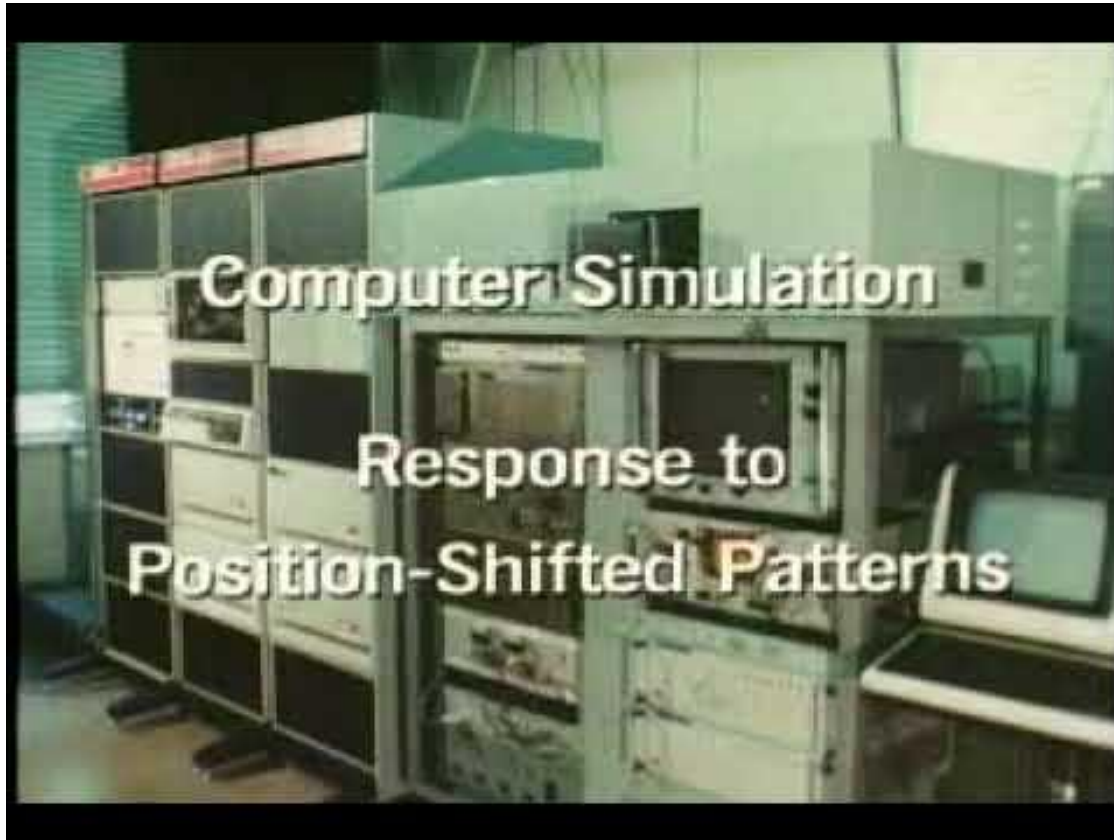
Visual Cortex Cell Recording



Neocognitron (1980): Parte I



Neocognitron (1980): Parte II



Backpropagation (1986)

Learning representations by back-propagating errors

David E. Rumelhart*, Geoffrey E. Hinton†
& Ronald J. Williams*

* Institute for Cognitive Science, C-015, University of California,
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,
Pittsburgh, Philadelphia 15213, USA

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector. As a result of the weight adjustments, internal 'hidden' units which are not part of the input or output come to represent important features of the task domain, and the regularities in the task are captured by the interactions of these units. The ability to create useful new features distinguishes back-propagation from earlier, simpler methods such as the perceptron-convergence procedure¹.

There have been many attempts to design self-organizing neural networks. The aim is to find a powerful synaptic modification rule that will allow an arbitrarily connected neural network to develop an internal structure that is appropriate for a particular task domain. The task is specified by giving the


more difficult when we introduce hidden units whose actual or desired states are not specified by the task. (In perceptrons, there are 'feature analysers' between the input and output that are not true hidden units because their input connections are fixed by hand, so their states are completely determined by the input vector: they do not learn representations.) The learning procedure must decide under what circumstances the hidden units should be active in order to help achieve the desired input-output behaviour. This amounts to deciding what these units should represent. We demonstrate that a general purpose and relatively simple procedure is powerful enough to construct appropriate internal representations.

The simplest form of the learning procedure is for layered networks which have a layer of input units at the bottom; any number of intermediate layers; and a layer of output units at the top. Connections within a layer or from higher to lower layers are forbidden, but connections can skip intermediate layers. An input vector is presented to the network by setting the states of the input units. Then the states of the units in each layer are determined by applying equations (1) and (2) to the connections coming from lower layers. All units within a layer have their states set in parallel, but different layers have their states set sequentially, starting at the bottom and working upwards until the states of the output units are determined.

The total input, x_j , to unit j is a linear function of the outputs, y_i , of the units that are connected to j and of the weights, w_{ji} , on these connections

$$x_j = \sum_i y_i w_{ji} \quad (1)$$

LeNet 5 (1988-1996)



Y. LeCun *et al.*, "Backpropagation Applied to Handwritten Zip Code Recognition," in *Neural Computation*, vol. 1, no. 4, pp. 541-551, Dec. 1989, doi: 10.1162/neco.1989.1.4.541.

LeNet 5 (1988-1996)



LeNet 5 (1988-1996)

Timeline


1989	Yann LeCun et al. proposed the original form of LeNet	LeCun, Y.; Boser, B.; Denker, J. S.; Henderson, D.; Howard, R. E.; Hubbard, W. & Jackel, L. D. (1989). Backpropagation applied to handwritten zip code recognition. Neural Computation, 1(4):541-551. ^[1]
1989	Yann LeCun proves that minimizing the number of free parameters in neural networks can enhance the generalization ability of neural networks.	LeCun, Y.(1989). Generalization and network design strategies. Technical Report CRG-TR-89-4, Department of Computer Science, University of Toronto. ^[2]
1990	Their paper describes the application of backpropagation networks in handwritten digit recognition once again	LeCun, Y.; Boser, B.; Denker, J. S.; Henderson, D.; Howard, R. E.; Hubbard, W. & Jackel, L. D. (1990). Handwritten digit recognition with a back-propagation network. Advances in Neural Information Processing Systems 2 (NIPS*89). ^[3]
1998	They reviewed various methods applied to handwritten character recognition and compared them with standard handwritten digit recognition benchmarks. The results show that convolutional neural networks outperform all other models.	LeCun, Y.; Bottou, L.; Bengio, Y. & Haffner, P. (1998). Gradient-based learning applied to document recognition.Proceedings of the IEEE. 86(11): 2278 - 2324. ^[4]

LeNet 5 (1988-1996)

Layer		Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	32x32	-	-	-
1	Convolution	6	28x28	5x5	1	tanh
2	Average Pooling	6	14x14	2x2	2	tanh
3	Convolution	16	10x10	5x5	1	tanh
4	Average Pooling	16	5x5	2x2	2	tanh
5	Convolution	120	1x1	5x5	1	tanh
6	FC	-	84	-	-	tanh
Output	FC	-	10	-	-	softmax

<https://www.datasciencecentral.com/profiles/blogs/lenet-5-a-classic-cnn-architecture#:~:text=The%20LeNet%2D5%20architecture%20consists,and%20finally%20a%20softmax%20classifier> .

LeNet 5 (1988-1996)



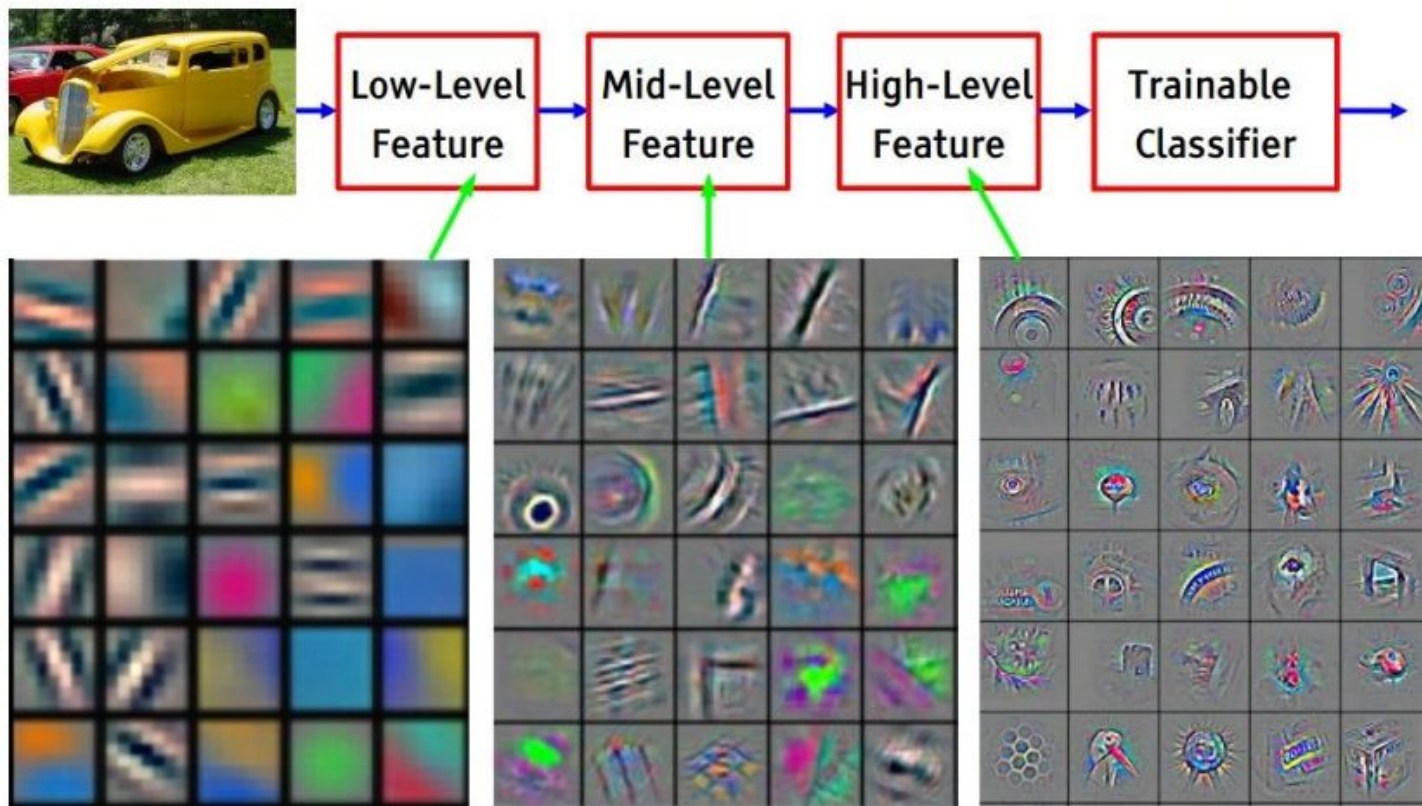
Y. Le Cun et. al, 1998

AlexNet (2012)


Layer		Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	227x227x3	-	-	-
1	Convolution	96	55 x 55 x 96	11x11	4	relu
	Max Pooling	96	27 x 27 x 96	3x3	2	relu
2	Convolution	256	27 x 27 x 256	5x5	1	relu
	Max Pooling	256	13 x 13 x 256	3x3	2	relu
3	Convolution	384	13 x 13 x 384	3x3	1	relu
4	Convolution	384	13 x 13 x 384	3x3	1	relu
5	Convolution	256	13 x 13 x 256	3x3	1	relu
	Max Pooling	256	6 x 6 x 256	3x3	2	relu
6	FC	-	9216	-	-	relu
7	FC	-	4096	-	-	relu
8	FC	-	4096	-	-	relu
Output	FC	-	1000	-	-	Softmax

Representaciones jerárquicas

Zeiler, M. D., & Fergus, R. (2014, September). Visualizing and understanding convolutional networks. In *European conference on computer vision* (pp. 818-833). Springer, Cham.



Brainscore



Integrative Benchmarking to Advance Neurally Mechanistic Models of Human Intelligence

Martin Schrimpf, Jonas Kubilius, Michael J. Lee, N. Apurva Ratan Murty, Robert Ajemian, James J. DiCarlo

Neuron

Volume 108 Issue 3 Pages 413-423
(November 2020)

DOI: 10.1016/j.neuron.2020.07.040

Brainscore



Brain-Score

Leaderboard

About

Compare

Participate

Rank	Model submitted by	average	V1 1 benchmark	V2 1 benchmark	V4 1 benchmark	IT 2 benchmarks	behavior	engineering 1 benchmark	Deng2009-top1 V1
1	CORnet-S Brain-Score Team	417	294	242	581	423	545	.747	.747
2	vgg-19 Brain-Score Team	408	347	341	610	248	494	.711	.711
3	resnet-50-robust Joel Dapello	408	378	365	537	243	515		
4	resnet-101_v1 Brain-Score Team	407	266	341	590	274	561	.764	.764
5	vgg-16 Brain-Score Team	406	355	336	620	259	461	.715	.715
6	resnet-152_v1 Brain-Score Team	405	282	338	598	277	533	.768	.768
7	resnet-101_v2 Brain-Score Team	404	274	332	599	263	555	.774	.774
8	resnet50-SIN_IN Brain-Score Team	404	282	324	599	276	541	.746	.746
9	densenet-169 Brain-Score Team	404	281	322	601	274	543	.759	.759
10	densenet-201 Brain-Score Team	402	277	325	599	273	537	.772	.772
11	resnet-50-pytorch Joel Dapello	399	289	317	600	259	528	.752	.752

<http://www.brain-score.org/#leaderboard>

Código en Keras

```
from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, Flatten

model = Sequential([
    Conv2D(16, 3, activation='relu', input_shape=(28,28,1)),
    MaxPool2D(),
    Conv2D(32, 3, activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(10, activation='softmax')
])
```